

Optimization

1. Unconstraint optimization problems

Andrew Lesniewski

Baruch College
New York

Fall 2018

Outline

- 1 Basics of unconstrained optimization
- 2 Nonlinear least squares problems
- 3 Quasi-Newton methods
- 4 Global optimization and differential evolution

Optimization problems

- Finance is all about *optimization*.
- Optimization helps to:
 - (i) maximize profits,
 - (ii) minimize costs,
 - (iii) control risk.
- Optimization problems facing financial professionals may include:
 - (i) Optimally allocate assets in a portfolio.
 - (ii) Build and interest rate curve based on benchmark rate observations.
 - (iii) Fit an implied volatility model to option market data.
 - (iv) Fit a time series model into observed data.
 - (v) Unwind a large position in a way that minimizes transaction costs.
 - (vi) Design an optimal hedging strategy of a portfolio of assets.
 - (vii) Optimally fit a predictive model into observed data.
 - (viii) Optimally exercise an American option.

Optimization problems

- Optimization is part of *applied mathematics*.
- That means that optimization problems are formulated in terms of properties of suitably defined functions.
- Depending on the context, solving an optimization problem draws on different disciplines of mathematics.
- A very limited number of optimization problems can be solved in “closed form”, i.e. a ready to use explicit formula. Typically, solving an optimization problem requires designing an *algorithm* and implementing it in *computer code*.
- Not all algorithms are equal: we strongly prefer algorithms with good performance characteristics.

Optimization problems

- Typically, *optimization problems* are formulated in terms of an *objective function* $f(x)$, $x \in \Omega \subset \mathbb{R}^n$ (depending on context, also known as a *loss function*, a *cost function*, etc.). The variables x are the *parameters* of the problem, and n is the *dimension* of the problem.
- The set $\Omega \subset \mathbb{R}^n$ defines constraints on the optimization problem, and its elements are called *feasible points*.
- More complex problems (such as multi-stage problems) may involve multiple objective functions.
- An optimization problem consists in finding a value x^* of the parameters such that $f(x) > f(x^*)$ for all values of x in a neighborhood of x^* . Such an x^* is called a *local minimizer* of $f(x)$.
- A *global minimizer* has the property that $f(x) > f(x^*)$ for *all* values of $x \in \Omega$.

Unconstrained optimization problems

- Finding global minimizers is typically much harder than finding local minimizer, as an objective function may have multiple local minima within Ω .
- In the case of a *convex* objective function, its local minimum is necessarily global, and there is no difference between the two problems. We will return to convex optimization later.
- A problem in which finding a maximum is required can be reduced to the case of a minimization problem by considering the objective function $\tilde{f}(x) = -f(x)$. In these lectures we will usually consider minimization problems.
- An optimization problem with $\Omega = \mathbb{R}^n$ are called *unconstrained*.

Necessary optimality conditions

- Let x^* be an unconstrained local minimizer of an objective function $f(x)$, $x \in \mathbb{R}^n$, and assume that $f(x)$ is continuously differentiable in a vicinity of x^* . Then:

- (i) Its gradient at x^* is zero,

$$\nabla f(x^*) = 0. \quad (1)$$

- (ii) If, additionally, $f(x)$ is twice continuously differentiable near x^* , then the Hessian at x^* is positive semidefinite:

$$\nabla^2 f(x^*) \geq 0. \quad (2)$$

- The proof is quite quick: Choose a direction $d \in \mathbb{R}^n$ and define a function $g(t) = f(x^* + td)$, $t \in \mathbb{R}$. Since x^* is a local minimizer,

$$\begin{aligned} 0 &\leq \lim_{t \rightarrow 0^+} \frac{g(t) - g(0)}{t} \\ &= g'(0) \\ &= \nabla f(x^*)^\top d. \end{aligned}$$

Necessary optimality conditions

- Since d is arbitrary, this is possible only if $\nabla f(x^*) = 0$.
- If $f(x)$ is twice continuously differentiable, then from Taylor's theorem:

$$\begin{aligned}g(t) &= g(0) + g'(0)t + \frac{1}{2} g''(0)t^2 + o(t^2) \\&= g(0) + \frac{1}{2} d^T \nabla^2 f(x^*) d t^2 + o(t^2).\end{aligned}$$

- Therefore,

$$\begin{aligned}0 &\leq \lim_{t \rightarrow 0} \frac{g(t) - g(0)}{t^2} \\&= d^T \nabla^2 f(x^*) d + \lim_{t \rightarrow 0} \frac{o(t^2)}{t^2} \\&= d^T \nabla^2 f(x^*) d,\end{aligned}$$

and we are done.

Sufficient optimality conditions

- Let the objective function $f(x)$, $x \in \mathbb{R}^n$, be twice continuously differentiable in an open set $O \subset \mathbb{R}^n$. Suppose that a point $x^* \in O$ satisfies the conditions

(i)

$$\nabla f(x^*) = 0. \quad (3)$$

- (ii) the Hessian at x^* is (strictly) positive definite:

$$\nabla^2 f(x^*) > 0. \quad (4)$$

- Then x^* is a strict local minimizer of $f(x)$.

Example: quadratic objective function

- Let $A \in \text{Mat}_n(\mathbb{R})$ be a positive definite matrix, $b \in \mathbb{R}^n$, and consider the objective function:

$$f(x) = \frac{1}{2} x^\top A x - x^\top b. \quad (5)$$

- The first order optimality condition reads

$$Ax^* = b,$$

i.e.

$$x^* = A^{-1}b. \quad (6)$$

- The Hessian

$$\nabla^2 f(x) = A \quad (7)$$

is positive definite (and constant). Therefore, (6) is a strict minimizer of (5).

Easy objective function: quadratic function


- An objective function that is easy to optimize is a quadratic function:

$$f(x, y) = (x - 1)^2 + (y - 1)^2. \quad (8)$$

- Inspection shows that it has a unique global minimum at $x = 1, y = 1$.
- This minimum is trivial to find analytically, and numerical algorithms usually have an easy time to find it.

Easy objective function: quadratic function

- The figure below shows a 3D graph of the quadratic function.



Hard objective function: Rosenbrock function


- A function, frequently used as a performance benchmark for testing optimization algorithms, is the *Rosenbrock function*:

$$f(x, y) = (x - 1)^2 + 100(y - x^2)^2. \quad (9)$$

- Inspection shows that it has a unique global minimum at $x = 1, y = 1$.
- This minimum, even though trivial to find analytically, lies inside a steep parabolic valley, and numerical algorithms struggle to find it.

Hard objective function: Rosenbrock function

- The figure below shows a 3D graph of the Rosenbrock function. Note the parabolic valley going through the point $(1, 1)$.



Sensitivity

- Suppose that we have a differentiable objective function that is parameterized by a set of parameters $\alpha \in \mathbb{R}^p$, i.e. the objective function is of the form $f(x, \alpha)$. Assume that for each α there is a local minimizer $x^*(\alpha)$, and that $x^*(\alpha)$ is a differentiable function of α .
- The first order optimality condition reads

$$\nabla_x f(x^*(\alpha), \alpha) = 0.$$

- Taking the gradient with respect to α and using the chain rule yields

$$\nabla_\alpha x^*(\alpha) \nabla_x^2 f(x^*(\alpha), \alpha) + \nabla_{x\alpha}^2 f(x^*(\alpha), \alpha) = 0,$$

where $\nabla_\alpha x(\alpha)$ is a $p \times n$ matrix of first derivatives $\partial x_i(\alpha)/\partial \alpha_k$. Assuming invertibility of the Hessian, this yields

$$\nabla_\alpha x^*(\alpha) = -\nabla_{x\alpha}^2 f(x^*(\alpha), \alpha) (\nabla_x^2 f(x^*(\alpha), \alpha))^{-1}. \quad (10)$$

Sensitivity

- This expresses the *sensitivity* of the minimizer to the parameter α .
- For example, if $f(x, \alpha) = f(x) - x^T \alpha$, $\alpha \in \mathbb{R}^n$, then

$$\nabla_{\alpha} x^*(\alpha) = (\nabla_x^2 f(x^*(\alpha)))^{-1}. \quad (11)$$

- Sensitivity analysis will be a helpful tool in studying constraint optimization problems.

NLS problems

- In a *nonlinear least square (NLS) problem* the objective function is of the form:

$$f(x) = \frac{1}{2} \sum_{i=1}^m \phi_i(x)^2, \quad x \in \mathbb{R}^n, \quad (12)$$

with given functions $\phi_i(x)$, and $m \geq n$. The factor $\frac{1}{2}$ is introduced for convenience.

- Introducing the column vector $\phi(x) = (\phi_1(x), \dots, \phi_m(x))^T$, we can conveniently write this objective function as

$$f(x) = \frac{1}{2} \phi(x)^T \phi(x). \quad (13)$$

- NLS problems arise in *data fitting*: We are given a set of data points (observations) $(T_1, S_1), \dots, (T_p, S_p)$ and a *data model* depending on a number of parameters x_1, \dots, x_k :

$$S = g(x, T). \quad (14)$$

The goal is to fit the data to the model.

- The model function $g(x, T)$ can be e.g. an explicitly defined function or a suitable spline function.

NLS problems

- Assume that

$$S_i = g(x, T_i) + \varepsilon_i, \quad (15)$$

where ε_i are residuals. One way to fit the model to the data is to minimize the sum of squares of the residuals, which leads to the following objective function:

$$f(x) = \frac{1}{2} \sum_{i=1}^m (g(x, T_i) - S_i)^2. \quad (16)$$

- In the simplest situation, x is two-dimensional, $x = (a, b) \in \mathbb{R}^2$, and the function $g(x, T)$ is linear:

$$\begin{aligned} S &= g(a, b, T) \\ &= a + bT. \end{aligned} \quad (17)$$

- This is the *linear* least square problem that can easily be solved in the closed form.

NLS problems

- The minimizer is given by the linear regression formulas:

$$\begin{aligned} b^* &= \frac{\sum_i (S_i - \bar{S})(T_i - \bar{T})}{\sum_i (T_i - \bar{T})^2}, \\ a^* &= \bar{S} - b^* \bar{T}, \end{aligned} \tag{18}$$

where $\bar{T} = \frac{1}{p} \sum_{i=1}^p T_i$, and $\bar{S} = \frac{1}{p} \sum_{i=1}^p S_i$.

- Except for the linear case, NLS problems cannot be solved in closed form. NLS problems tend to be hard (think about it: the Rosenbrock function can be viewed as an NLS problem), and require refined numerical algorithms.

NLS problems

- In finance, NLS problems arise when we face the task of interpolating a cross section of market data. This amounts to the problem of *curve building* in a variety of contexts. Here are a few important examples.
- Given a snapshot of benchmark swap rates of given tenors, build a smooth interpolating curve that allows us to compute the swap rate of any tenor.
- Given a snapshot of the US Treasuries, fit a smooth curve that allows to carry out rich-cheap analysis.
- Given a snapshot of benchmark credit spreads, build a smooth interpolating curve that allows us to compute the spread of any maturity.
- Given a snapshot of option prices of the same expiration (and underlying) but varying strikes, build an interpolation curve that allows us to predict volatility smile for any strike.

NLS problems

- We will now develop an number of algorithms for solving NLS problems. Each of these algorithms is *recursive*, meaning that it starts with an initial guess for x^* and continues until an exit criterion is met.
- Before we move on, notice that the gradient of the objective function (12) (or (13)) is given by

$$\begin{aligned}\nabla f(x) &= \sum_{i=1}^m \phi_i(x) \nabla \phi_i(x) \\ &= J(x)^T \phi(x),\end{aligned}\tag{19}$$

where $J(x)_{ij} = \nabla_j \phi_i(x)$ is the $m \times n$ Jacobian matrix.

- Likewise, the Hessian of $f(x)$ is given by

$$\begin{aligned}\nabla^2 f(x) &= \sum_{i=1}^m (\nabla \phi_i(x) \nabla \phi_i(x)^T + \phi_i(x) \nabla^2 \phi_i(x)) \\ &= J(x)^T J(x) + \sum_{i=1}^m \phi_i(x) \nabla^2 \phi_i(x).\end{aligned}\tag{20}$$

NLS problems

- As discussed earlier, finding a local minimizer x^* of $f(x)$ requires solving the equation

$$\nabla f(x^*) = 0. \quad (21)$$

- If x is close to x^* , we can use Taylor's expansion to approximate

$$\begin{aligned} 0 &= \nabla f(x^*) \\ &\approx \nabla f(x) + \nabla^2 f(x)(x^* - x). \end{aligned}$$

- Assuming that the matrix $\nabla^2 f(x)$ is nonsingular for all x , this suggests the following *Newton method* update rule:

(i) Start with an initial guess x_0 for x^* .

(ii) For $k = 1, 2, \dots$,

$$x_{k+1} = x_k - \nabla^2 f(x_k)^{-1} \nabla f(x_k). \quad (22)$$

- This observation is the starting point of several powerful algorithms for solving the NLS problem.

- In some ways, the scheme above is very attractive: it can be proved that if x_0 is sufficiently close to the local minimizer x^* , then the sequence x_k converges rapidly to x^* , namely quadratically in the error $\|x_{k+1} - x_k\|$.
- Notice, however, some of the practical challenges that it faces:
 - (i) Finding an appropriate initial guess may be difficult.
 - (ii) Expressions (19) and (20) for the gradient and Hessian of $f(x)$, respectively, require computing the first and second derivatives of the functions $\phi_i(x)$. Preferably, these derivatives should be calculated in closed form. Typically, even in the case of moderate dimensionality n of the problem, this is infeasible. It is especially true for the second derivatives.
 - (iii) One may try to calculate the derivatives numerically, as finite differences. In the case of a moderate / large dimensionality of the problem, this is computationally prohibitively expensive, and the algorithm is inefficient.
 - (iv) There is no guarantee that the Hessian in (22) is invertible. Even if it is, there is a cost associated with this operation at each iteration step.

Gradient descent method

- The *gradient descent method* is the simplest approximation to (22), in which assume that the Hessian is a multiple of the identity matrix.
- The scheme can thus be formulated as
 - (i) Start with an initial guess x_0 for x^* .
 - (ii) For $k = 1, 2, \dots$,

$$x_{k+1} = x_k - \alpha \nabla f(x_k). \quad (23)$$

Here the parameter $\alpha > 0$ is called the *step size* or *learning rate*.

- The value of α can be either determined by trial and error or by means of a *line search* algorithm. In the latter approach, in each step, we numerically search for a value α_k satisfying

$$\alpha_k = \arg \min_{\alpha} f(x_k - \alpha \nabla f(x_k)). \quad (24)$$

- Gradient descent methods are easy to implement, but their performance tends to be mixed. Convergence (if any) is linear and slow, and it requires many iterations.
- Among the descent methods, the gradient descent method is locally the best, as it moves the next iteration in the direction of the steepest descent of $f(x)$.

Gauss-Newton method

- The *Gauss-Newton method* is based on the assumption that the function $\phi_i(x)$ can be locally well approximated by linear functions:

$$\phi(x + h) \approx \phi(x) + J(x)h. \quad (25)$$

- This implies that

$$\begin{aligned} f(x + h) &\approx f(x) + \nabla f(x)h + \frac{1}{2} h^T \nabla^2 f(x)h \\ &= f(x) + h^T J(x)^T \phi(x) + \frac{1}{2} h^T J(x)^T J(x)h, \end{aligned} \quad (26)$$

where we have used expressions (19) and (20) for the gradient and Hessian of $f(x)$, respectively.

- Notice the absence of the second derivatives in (20) from the term quadratic in h . This is a consequence of the linear approximation assumption. This is a boon to the computational efficiency of algorithm.

Gauss-Newton method

- Effectively, we thus approximate the Hessian of $f(x)$ as follows:

$$\nabla^2 f(x) \approx J(x)^\top J(x). \quad (27)$$

This approximation guarantees that Hessian is always positive semidefinite. Also, it turns out that dropping the second derivatives does not, in practice, affect the convergence of the process.

- We now choose the optimal h by minimizing the quadratic function in h on the right hand side of (26). This yields the condition

$$J(x)^\top J(x)h = -J(x)^\top \phi(x). \quad (28)$$

- The function $f(x)$ descends along the direction of h as

$$\begin{aligned} h^\top \nabla f(x) &= J(x)^\top h \\ &= -h^\top J(x)^\top J(x)h \\ &\leq 0. \end{aligned}$$

Gauss-Newton method

- The Gauss-Newton scheme can be thus formulate as follows. We start at x_0 and iterate:

$$x_{k+1} = x_k - (J(x_k)^\top J(x_k))^{-1} J(x_k)^\top \phi(x_k). \quad (29)$$

The key step of the method is solving equation (31), which amounts to inverting the matrix $J(x_k)^\top J(x_k)$. In case if $J(x)$ is poorly conditioned, this inverse may not exist.

- It may be beneficial to refine the method by introducing the learning rate α ,

$$x_{k+1} = x_k - \alpha (J(x_k)^\top J(x_k))^{-1} J(x_k)^\top \phi(x_k), \quad (30)$$

which can be either set *a priori* or determined at each step by conducting a line search. Under suitable technical assumptions, this method is guaranteed to converge.

- Unlike the Newton method, the Gauss-Newton method typically converges linearly in h .

Levenberg-Marquardt method

- The *Levenberg-Marquardt method* is a clever hybrid of the gradient descent and Gauss-Newton methods.
- Namely, we introduce a *damping parameter* $\mu > 0$, and determine the step size by solving the equation

$$(J(x)^T J(x) + \mu I) h = -J(x)^T \phi(x). \quad (31)$$

- Notice that the presence of μ has the following effects on the problem:
 - (i) The matrix $J(x)^T J(x) + \mu I$ is strictly positive definite, and the existence of its inverse is guaranteed.
 - (ii) For small values of μ , the iteration step is similar to the iteration step in the Gauss-Newton method.
 - (iii) For large values of μ , the iteration step is similar to the iteration step in the gradient descent method.
- The value of μ will be updated throughout the iteration process, and it renders line search unnecessary.

Levenberg-Marquardt method

- In order to initialize μ , one usually chooses it be of the order of magnitude of maximum diagonal term of the matrix $J(x_0)^T J(x_0)$,

$$\mu_0 = \beta \max_i (J(x_0)^T J(x_0))_{ii}, \quad (32)$$

where $\beta \approx 10^{-3}$ is determined empirically.

- The update process is determined bu the *gain ratio*, defined as the ratio of the projected change in $f(x)$ to the change in $f(x)$ projected by the quadratic approximation (26):

$$\rho = - \frac{f(x_k) - f(x_k + h)}{h^T J(x_k)^T \phi(x_k) + \frac{1}{2} h^T J(x_k)^T J(x_k) h}, \quad (33)$$

which can also be expressed as

$$\rho = 2 \frac{f(x_k) - f(x_k + h)}{h^T (\mu h - J(x_k)^T \phi(x_k))}. \quad (34)$$

Levenberg-Marquardt method

- Note that the denominator in (34) is always positive (why?)
- Large value of ρ indicates that the quadratic approximation to $f(x)$ is accurate. We should decrease the value of μ in order to make the process more Gauss-Newton like, and accelerate the convergence.
- On the other hand, a small value of ρ (also negative) indicates that the approximation is poor and we are on the wrong path. In this case, the value of μ should be increased to make the process more gradient descent like, and also decrease the step size.
- The updates of μ are done in terms of a scaling factor λ (frequently chosen as $\lambda = 2$): if the gain ratio is positive, the value of μ is reduced by a factor depending on ρ , otherwise it is scaled up according to $\mu \rightarrow \lambda\mu$.

Levenberg-Marquardt method

- An important element of a search algorithm is a set of *exit criteria*, which stop the algorithm once it has converged or it failed to converge.
 - (i) Ideally, the search should stop if $\nabla f(x^*) = J(x^*)^T \phi(x^*) = 0$. We can implement it numerically by choosing a small $\varepsilon_1 > 0$, and require that all components of the vector $\nabla f(x_k)$ have absolute values less than ε_1 . This is equivalent to stating that

$$\|J(x_k)^T \phi(x_k)\|_\infty \leq \varepsilon_1. \quad (35)$$

Here, the L^∞ -norm $\|X\|_\infty$ of a vector X is defined as $\max_j |X_j|$.

- (ii) Another criterion guarantees exit, if the changes in x_k are slow. We choose another small number ε_2 and verify if

$$\|x_{k+1} - x_k\| \leq \varepsilon_2(\|x_k\| + \varepsilon_2). \quad (36)$$

- (iii) Lastly, in order to avoid running an infinite loop, one has to choose the maximum number of iterations k_{max}

Levenberg-Marquardt method

- Let us now summarize the Levenberg-Marquardt method in an algorithmic form:

choose x_0 , initialize μ according to (32), and $\lambda = 2$

done = $(\|J(x_0)^\top \phi(x_0)\|_\infty \leq \varepsilon_1)$

while (*not done* and $k \leq k_{max}$)

$h = -(J(x_k)^\top J(x_k) + \mu I)^{-1} J(x_k) \phi(x_k)$

if $\|h\| \leq \varepsilon_2(\|x_k\| + \varepsilon_2)$

done = *true*

else

calculate ρ according to (33)

if $\rho > 0$

$x_{k+1} = x_k + h$

done = $(\|J(x_{k+1})^\top \phi(x_{k+1})\|_\infty \leq \varepsilon_1)$

update $\mu \rightarrow \mu \max(1/3, 1 - (2\rho - 1)^3)$, and $\lambda = 2$

else

update $\mu \rightarrow \lambda\mu$, and $\lambda \rightarrow 2\lambda$

Levenberg-Marquardt method

- The Levenberg-Marquardt method is implemented in `scipy`.

Data fitting without NLS

- Data fitting problems in finance often lead to optimization problems which are not of the NLS type.
- For example, in models requiring maximum likelihood estimation (MLE) from historical data, the objective function is typically not a sum squares of the residuals between the data and the model. Instead, it is the probability of the observed data given given the (unknown) parameters of the probability distribution.
- For example, many problems lead to the task of minimizing the Kullback-Leibler divergence between the observed frequencies p and a probability model $q(x)$,

$$KL(p||q(x)) = \sum_{i=1}^m p_i \log \frac{p_i}{q_i(x)} . \quad (37)$$

This objective function clearly is not of the NLS form.

- Quasi-Newton methods have been designed to problems like this. The most popular among them is the *BGFS method* (for Broyden, Fletcher, Goldfarb, and Shanno) and its variations.

BFGS method

- Our goal is to find a local minimizer of an objective function $f(x)$, not necessary expressed as a sum of squares. We initialize the search with x_0 and iterate.
- At the $(k+1)$ -th step of iteration we approximate the value of $f(x_{k+1})$ by an expression quadratic in the direction of the step from x_k . Namely, for x in the vicinity of x_k , we can write

$$f(x) \approx f(x_k) + \nabla f(x_k)^\top h + \frac{1}{2} h^\top B_k h, \quad (38)$$

where B_k is a positive definite matrix.

- Naively, B_k should be the Hessian at x_k , but, as we know, this leads to serious performance issues, and the challenge is to choose B_k smartly.
- The minimizer of (38) is

$$h = -B_k^{-1} \nabla f(x_k), \quad (39)$$

and x_{k+1} is given by

$$x_{k+1} = x_k + \alpha_k h. \quad (40)$$

BFGS method

- The value of the multiplier α_k is determined from the line search:

$$\alpha_k = \arg \min_{\alpha} f(x_k - \alpha \nabla f(x_k)). \quad (41)$$

- In the BFGS method, the matrix B_k is constructed iteratively as follows. Given B_k , we require that the updated value B_{k+1} satisfies the *secant equation*

$$B_{k+1}(x_{k+1} - x_k) = \nabla f(x_{k+1}) - \nabla f(x_k) \quad (42)$$

(remember B_k is supposed to mimic the Hessian!).

- This can be written as

$$B_{k+1}s_k = y_k, \quad (43)$$

where $s_k = x_{k+1} - x_k$ and $y_k = \nabla f(x_{k+1}) - \nabla f(x_k)$.

BFGS method

- Premultiplying (43) by s_k^\top , we see that the following *curvature condition* has to be satisfied:

$$s_k^\top y_k > 0. \quad (44)$$

If the curvature condition is satisfied, the secant equation (43) has a unique solution.

- For non-convex functions, the curvature condition is not always satisfied. In such a case, it has to be enforced by means of choosing an appropriate step size α_k .
- We will now assume that the updated B_{k+1} equals to B_k plus two positive definite rank 1 matrices:

$$B_{k+1} = B_k + auu^\top + bvv^\top, \quad (45)$$

where $a, b \in \mathbb{R}$, and $u, v \in \mathbb{R}^n$.

- How do we choose the parameters above?

BFGS method

- This in itself is an optimization in which we minimize the Frobenius norm between the symmetric matrices B_{k+1} and B_k under the constraint that the secant equation holds. The solution turns out to be

$$\begin{aligned}u &= y_k, \\v &= B_k s_k \\a &= \frac{1}{s_k^\top y_k}, \\b &= -\frac{1}{s_k^\top B_k s_k},\end{aligned}$$

which leads to the following update equation for B_{k+1} :

$$B_{k+1} = B_k - \frac{B_k s_k s_k^\top B_k}{s_k^\top B_k s_k} + \frac{y_k y_k^\top}{s_k^\top y_k}. \quad (46)$$

BFGS method

- We are now almost ready, except for the choice of the initial matrix B_0 . Unfortunately, there is no standard recipe for the choice of B_0 .
- One possible choice is the Hessian of $f(x)$ calculated at the initial guess x_0 (either in closed form, or using finite differences.)
- Another choice is a multiple of the identity matrix, $B_0 = aI$, where a is user determined.
- For the exit criterion we use $\|\nabla f(x_k)\|_\infty < \varepsilon$, where ε is user determined.

BFGS method

- Let us now summarize the BFGS method in an algorithmic form:

initialize x_0 and B_0

while $(\|\nabla f(x_k)\|_\infty > \varepsilon)$ and $k \leq k_{max}$

 compute the search direction (39)

 carry out the line search (41)

 set $s_k = \alpha_k h$, update $x_{k+1} = x_k + s_k$, and calculate y_k

 update B_k to B_{k+1} according to equation (46)

BFGS method

- The algorithm above requires inverting the matrix B_k at each step of the iteration.
- There is a version of the BFGS method, which operates in terms of the inverses $H_k = B_k^{-1}$ only.
- Namely, one can verify (with some pain) that the update equation (46) can be formulated in terms of the inverses as follows:

$$H_{k+1} = \left(I - \frac{s_k y_k^\top}{s_k^\top y_k} \right) H_k \left(I - \frac{y_k s_k^\top}{s_k^\top y_k} \right) + \frac{s_k s_k^\top}{s_k^\top y_k}. \quad (47)$$

- The BFGS method is implemented in the Python package `scipy`.

Global optimization

- In many applications, the objective function has multiple local minima, or at least the existence of a unique global minimum cannot be guaranteed.
- Furthermore, objective functions can be non-differentiable, noisy, exhibit flat directions, etc.
- In situations like this, analytic methods or Newton-like methods are not applicable.
- A number of approaches for solving such optimization problems have been proposed. We will focus on a particularly neat stochastic search method, namely *differential evolution*.

Hard objective function: Ackley function


- Another function, frequently used as a performance benchmark for testing optimization algorithms, is the *Ackley function*:

$$f(x, y) = 20 + e - 20 \exp \left(-0.2 \sqrt{0.5(x^2 + y^2)} \right) - \exp \left(0.5 (\cos(2\pi x) + \cos(2\pi y)) \right). \quad (48)$$

- Inspection shows that it has a unique global minimum at $x = 0, y = 0$.
- This minimum, even though trivial to find analytically, is hidden among plenty of local minima and maxima, which makes Newton-type methods impractical.

Hard objective function: Ackley function

- The figure below shows a 3D graph of the Ackley function. Note the multitude of local minima and maxima.



Differential evolution

- As before, let $f(x)$, $x \in \mathbb{R}^n$ be an objective function. We do not require $f(x)$ to be differentiable.
- Each of the components of x is restricted to a finite interval:

$$L_j \leq x_j \leq U_j. \quad (49)$$

The limits are chosen by the user. Tighter limits lead to a better performance of the algorithm.

- We select the population size $N \geq 4$, and randomly generate the initial population: $x_{1,0}, \dots, x_{N,0}$.
- We will evolve this population through a number of generations G according to the rules motivated by genetics.

Differential evolution

- *Mutation*. We choose a constant *mutation factor* $F \in [0, 2]$, determined by the user.
- For a given vector $x_{i,G}$ (we start with $G = 0$), we randomly select three vectors $x_{r_1,G}, x_{r_2,G}, x_{r_3,G}$, such that the four indices i, r_1, r_2, r_3 are all distinct.
- We form the *donor (mutant) vectors*:

$$v_{i,G+1} = x_{r_1,G} + F(x_{r_2,G} - x_{r_3,G}), \quad (50)$$

for each element of the population $i = 1, \dots, N$.

Differential evolution

- *Crossover*. We choose a user determined *crossover constant* $CR \in [0, 1]$.
- For each $i = 1, \dots, N$ we form a *trial vector* $u_{i,G+1}$ constructed as follows.
 - (i) Generate random numbers $u_j \sim U(0, 1), j = 1 \dots, n$.
 - (ii) Generate a random integer r from the range $1, \dots, n$.
 - (iii) For $j = 1 \dots, n$, set

$$u_{j,i,G+1} = \begin{cases} v_{j,i,G+1}, & \text{if } u_j \leq CR \text{ or } j = r, \\ x_{j,i,G}, & \text{if } u_j > CR \text{ and } j \neq r. \end{cases} \quad (51)$$

Differential evolution

- *Selection.* The target vector $x_{i,G}$ is now compared with the trial vector $u_{i,G+1}$, and the one with the lower value of the objective function is passed on to the next generation.
- Namely, for $i = 1, \dots, N$,

$$x_{i,G+1} = \begin{cases} u_{i,G+1}, & \text{if } f(u_{i,G+1}) < f(x_{i,G}) \\ x_{i,G} & \text{otherwise.} \end{cases} \quad (52)$$

Differential evolution

- Choosing the right parameters for the differential evolution requires trial and error. Usually, $F = 0.5$, and $CR = 0.1$ work well. Population size N should be a function of the dimensionality of the problem. A useful rule of thumb is $N = 5n$.
- Implementation of differential evolution requires exit criteria imposed by the user. One of them is the maximum number of generations G_{max} .
- Another criterion is the average rate of improvement of $\min_i f(x_{i,G})$ from generation to generation. Once this rate falls below a user chosen level ε , the algorithm should stop.

References



Madsen, K., Nielsen, K. B., and Tingleff, O.: *Methods for Non-Linear Least Squares Problems* Technical University Denmark (2014).



Nocedal, J. and Wright, S. J.: *Numerical Optimization*, Springer (2006).



Storn, R., and Price, K.: Differential evolution - a simple and efficient heuristic for global optimization over continuous spaces, *Journal of Global Optimization*, **11**, 341–359 (1997).